



Context-Length Inference From Cambricon-LLM Figure 9 Fit



Zoom Around the Valid Fit Window



Inference Summary

Default context **1,000 tokens**

Guardrail window **970-1,040 tokens**

Best max-error fit **990 tokens**

Best RMSE fit **1,023 tokens**

Default mean abs error **8.341%**

Default max error **14.618%**

Interpretation: the public Figure 9 fit supports an inferred context length near 1K; the paper text is not treated as explicitly specifying this value.