


The public Figure 9 fit supports an inferred 1K decode context

a

Context-length sweep against all 21 public Figure 9 points



b

Fit summary

1000 tokens

Default context

975-1038

Guardrail window

1005 tokens

Best max-error fit

1030 tokens

Best RMSE fit

8.356%

Default MAE

14.508%

Default max error